

Introduction to R

Tidyverse-I

Sumit Mishra

Krea University | WSDS002

Contents

Tibbles	1
Creation	2
Importing a dataset	2
dplyr	3
Intro	3
Manipulating variables: mutate()	3
Pipes	4
%>% and dplyr	5
select()	5
relocate()	7
summarize()	8
summarize() and group_by()	8
filter()	9
arrange()	9
slice()	10
distinct()	10
Chain Operation Revisited	10
Done for the day	11

Tibbles

tibble is the modern version of R's data.frame.

Example starwars dataset that is a built-in tibble.

```
## [1] "tbl_df"      "tbl"        "data.frame"
```

name	height	mass	sex	gender	homeworld
Luke Skywalker	172	77	male	masculine	Tatooine
C-3PO	167	75	none	masculine	Tatooine
R2-D2	96	32	none	masculine	Naboo
Darth Vader	202	136	male	masculine	Tatooine
Leia Organa	150	49	female	feminine	Alderaan
Owen Lars	178	120	male	masculine	Tatooine

Creation

The tibble function creates a data.frame like-object. You'll generally define tibble by passing the column names and values for the columns.

```
album <- c("Please Please Me", "Rubber Soul", "Magical Mystery Tour")
year <- c(1963, 1965, 1967)
num.tracks <- c(14, 14, 11)
beatles.catalog <- tibble(album, year, num.tracks)
```

You can also create tibbles from other objects (e.g., matrices) using the function `as_tibble()`¹

Importing a dataset

- `read_csv`: reads csv file into R as a tibble.

```
g3 <- read_csv("../data/gdp-growth.csv")
```

```
## Parsed with column specification:
## cols(
##   .default = col_double(),
##   `Country Name` = col_character(),
##   `Country Code` = col_character(),
##   `Series Name` = col_character()
## )
## See spec(...) for full column specifications.
```

```
head(g3)
```

```
## # A tibble: 6 x 23
##   `Country Name` `Country Code` `Series Name` YR1999 YR2000 YR2001 YR2002 YR2003
##   <chr>         <chr>         <chr>      <dbl>  <dbl>  <dbl>  <dbl>  <dbl>
## 1 Afghanistan  AFG           GDP growth (~ NA      NA      NA      NA      8.83
## 2 Albania      ALB           GDP growth (~ 12.9    6.95    8.29    4.54    5.53
## 3 Algeria      DZA           GDP growth (~ 3.20    3.82    3.01    5.61    7.20
## 4 American Samoa ASM           GDP growth (~ NA      NA      NA      NA      0.814
## 5 Andorra      AND           GDP growth (~ 4.10    3.53    4.55    6.47    12.2
## 6 Angola       AGO           GDP growth (~ 2.18    3.05    4.21    13.7    2.99
## # ... with 15 more variables: YR2004 <dbl>, YR2005 <dbl>, YR2006 <dbl>,
## #   YR2007 <dbl>, YR2008 <dbl>, YR2009 <dbl>, YR2010 <dbl>, YR2011 <dbl>,
## #   YR2012 <dbl>, YR2013 <dbl>, YR2014 <dbl>, YR2015 <dbl>, YR2016 <dbl>,
## #   YR2017 <dbl>, YR2018 <dbl>
```

- `read_xls`: reads an excel file. You will need to have the package `readxl` installed and loaded.

```
col <- read_xls("../data/COL.xls")
```

```
## # A tibble: 65 x 13
##   year population y_pop y_pop_us_100 y_wkr y_wkr_us_100 growth_100 i_y g_y
##   <dbl>      <dbl> <dbl>      <dbl> <dbl>      <dbl> <chr>      <dbl> <dbl>
## 1 1950      11.9  3143      21.5  9820      27.1 NaN      25.8  6.89
## 2 1951      12.3  3092      20.0  9786      26.1 -1.6431  25.1  7.74
## 3 1952      12.7  3196      20.4 10247      26.7 3.3111000~ 26.2  7.76
## 4 1953      13.1  3364      20.8 10926      27.6 5.1102999~ 27.5  8.68
## 5 1954      13.5  3571      22.6 11751      29.3 5.9764999~ 30.7  8.22
## 6 1955      14.0  3591      21.5 11963      28.5 0.5639999~ 31.1  8.43
```

¹Or just plain, old `tibble()`.

```
## 7 1956      14.4 3557      21.2 11978      28.6 -0.951500~ 31.0 8.03
## 8 1957      14.9 3436      20.4 11695      27.4 -3.469800~ 28.6 7.22
## 9 1958      15.4 3319      20.2 11423      26.6 -3.444100~ 25.0 7.60
## 10 1959      15.9 3405      19.7 11846      26.3 2.5419999~ 25.0 7.22
## # ... with 55 more rows, and 4 more variables: ed_att <chr>, nx_y <dbl>,
## #   x_m_y <dbl>, ex_rate <dbl>
```

Year	Population	Y/Pop	Y/Pop(us=100)	Y/Wkr
1950	11.9333	3143	21.5011	9820
1951	12.3176	3092	20.0197	9786
1952	12.7148	3196	20.3629	10247
1953	13.1254	3364	20.8274	10926
1954	13.5499	3571	22.5808	11751

- R is pretty versatile in getting all kinds of datasets. For example, you can use `fromJSON()` to read json files from the internet.

```
cov19district <- jsonlite::fromJSON("https://api.covid19india.org/state_district_wise.json", flatten=T)
```

dplyr

Intro

It's a package. dplyr is not installed by default, so you'll need to install it².

dplyr is part of the [tidyverse](#), and it follows a grammar-based approach to programming/data work.

- data compose the subjects of your stories
- dplyr provides the *verbs* (action words): `filter()`, `mutate()`, `select()`, `group_by()`, `summarize()`, `arrange()`

Manipulating variables: mutate()

dplyr streamlines adding/manipulating variables in your data frame.

Function: `mutate(.data, ...)`

- **Required argument:** `.data`, an existing data frame
- **Additional arguments:** Names and values of the new variables
- **Output:** An updated data frame

Example Take the data frame

```
df <- tibble(x = seq(2,22, length.out = 6),
            y = sample(1:20, 6))
```

`mutate()` allows us to create many new variables with one call.

²or just `p_load(dplyr)` after loading `pacman`

Code:

```
mutate(.data = df,
  xy = x * y,
  y2 = y^2,
  y_x = round(y/x),
  is_y_min = y == min(y)
)
```

Output:

```
## # A tibble: 6 x 6
##       x     y   xy   y2  y_x is_y_min
##   <dbl> <int> <dbl> <dbl> <dbl> <lgl>
## 1     2    20    40   400    10 FALSE
## 2     6     1     6     1     0  TRUE
## 3    10    10   100   100     1 FALSE
## 4    14     5    70    25     0 FALSE
## 5    18    14   252   196     1 FALSE
## 6    22    16   352   256     1 FALSE
```

Please note that `mutate()` returns the original *and* new columns.

Pipes

Before we go further, let's take a detour to learn about an important operator in tidyverse: pipe `%>%`. A *pipe* in programming allows you to take the output of one function and plug it into another function as an argument/input.

R's pipe specifically plugs the returned object to the left of the pipe into the first argument of the function on the right for the pipe, e.g.,

```
seq(2,22,length.out = 6) %>% mean() %>% round()
```

```
## [1] 12
```

Pipes help avoid lots of nested functions, and increase the readability of our code.

Example We will randomly pick six numbers between 5 and 30, compute their average, and round off the average. Remember the workflow.

Numbers → Sample → Average → Round off

```
# Save each intermediate step
numbers <- 5:30
our_sample <- sample(numbers,6)
ave.num <- round(mean(our_sample))
# Lots of nesting
ave.num <- round(mean(sample(5:30,6)))
print(ave.num)
```

```
## [1] 16
```

```
# Piping
ave.num <- 5:30 %>% sample(6) %>% mean() %>% round()
print(ave.num)
```

```
## [1] 13
```

By default, R pipes the output from the LHS of the pipe into the first argument of the function on the RHS of the pipe.

E.g., `x %>% rep(3)` is equivalent to `rep(x, size = 3)`.

If you want to pipe output into a different argument, you use a period (`.`).

Example Suppose that you have a vector `x` of length 100, and you want to generate a sample `y` of size 10. You can achieve this using pipe in the following different ways.

```
x <- rnorm(100)
```

- Option 1

```
y <- x %>% sample(.,10,replace=F)
print(y)
```

```
## [1] 0.9204039 0.6016666 -0.6007949 2.2084500 -0.1421788 -0.4025734
## [7] -0.1024816 0.6042069 2.4706415 0.8028889
```

- Option 2

```
y <- 10 %>% sample(x,., replace = F)
print(y)
```

```
## [1] 0.9204039 -0.1383846 -0.6007949 -0.9526688 0.4776354 -0.3674844
## [7] -0.1065338 0.7058072 -1.1933317 1.0167047
```

- Option 3

```
y <- (replace = F) %>% sample(x,10,.)
print(y)
```

```
## [1] -1.53744735 0.23736545 0.39136145 0.62136298 -0.28991560 -0.29567242
## [7] 0.57790658 -0.75504410 -0.28112169 -0.08260369
```

%>% and dplyr

Each dplyr function begins with a `.data` argument so that you can easily pipe in data frames (recall: `mutate(.data, ...)`).

The common workflow in dplyr will look something like

```
new_df <- old_df %>% mutate(cool stuff here)
```

which takes `old_df`, does some cool stuff with `mutate()`, and then saves the output of `mutate()` as `new_df`. Saving as a new (or replace the old) data frame helps you use the newly created columns.

Example

Without pipe:

```
new_df <-
  mutate(.data = df,
    xy = x * y,
    y2 = y^2,
    y_x = round(y/x),
    is_y_min = y == min(y)
  )
```

Pipe:

```
# show output
new_df <- df %>%
  mutate(xy = x * y,
    y2 = y^2,
    y_x = round(y/x),
    is_y_min = y == min(y)
  )
```

select()

Just as `filter()` outputs row-based subsets of your tibble, `select()` grabs **column-based subsets**.

You can select columns using their **names** `new_df %>% select(xy, x)`

or you can select columns using **helper functions** `new_df %>% select(starts_with("x"))`

You can also choose to drop a column by prefixing the name of the column by hyphen (`-`).

```
beatles.catalog %>% select(-num.tracks)
```

```
## # A tibble: 3 x 2
##   album          year
##   <chr>         <dbl>
```

```
## 1 Please Please Me      1963
## 2 Rubber Soul           1965
## 3 Magical Mystery Tour  1967
```

Renaming variables can also be done using `select()`. The syntax will be simple: `select(NEW NAME = OLD NAME)`. Example:

```
starwars %>%
  select(alias=name, crib=homeworld, sex=gender)
```

```
## # A tibble: 87 x 3
##   alias      crib      sex
##   <chr>      <chr>   <chr>
## 1 Luke Skywalker   Tatooine masculine
## 2 C-3PO            Tatooine masculine
## 3 R2-D2            Naboo    masculine
## 4 Darth Vader      Tatooine masculine
## 5 Leia Organa      Alderaan feminine
## 6 Owen Lars        Tatooine masculine
## 7 Beru Whitesun lars Tatooine feminine
## 8 R5-D4            Tatooine masculine
## 9 Biggs Darklighter Tatooine masculine
## 10 Obi-Wan Kenobi   Stewjon  masculine
## # ... with 77 more rows
```

Select helpers

- `starts_with()`: Starts with a prefix

Example Select country names and GDP variables from `g3`.

```
g3 %>% select(`Country Name`, starts_with("YR"))
```

```
## # A tibble: 264 x 21
##   `Country Name` YR1999 YR2000 YR2001 YR2002 YR2003 YR2004 YR2005 YR2006
##   <chr>          <dbl>  <dbl>  <dbl>  <dbl>  <dbl>  <dbl>  <dbl>  <dbl>
## 1 Afghanistan    NA     NA     NA     NA     8.83   1.41  11.2   5.36
## 2 Albania         12.9   6.95   8.29   4.54   5.53   5.51   5.53   5.90
## 3 Algeria         3.20   3.82   3.01   5.61   7.20   4.30   5.91   1.68
## 4 American Samoa  NA     NA     NA     NA     0.814  0.538 -0.402 -4.17
## 5 Andorra         4.10   3.53   4.55   6.47  12.2   7.65   7.40   4.54
## 6 Angola          2.18   3.05   4.21  13.7   2.99  11.0   15.0   11.5
## 7 Antigua and B~  3.71   6.69  -4.95   1.02   6.06   5.74   6.41  12.7
## 8 Arab World      1.80   5.48   1.61   0.586  5.32   9.34   5.72   6.50
## 9 Argentina      -3.39 -0.789 -4.41 -10.9   8.84   9.03   8.85   8.05
## 10 Armenia        3.30   5.9    9.56  13.2   14.0   10.5   13.9   13.2
## # ... with 254 more rows, and 12 more variables: YR2007 <dbl>, YR2008 <dbl>,
## #   YR2009 <dbl>, YR2010 <dbl>, YR2011 <dbl>, YR2012 <dbl>, YR2013 <dbl>,
## #   YR2014 <dbl>, YR2015 <dbl>, YR2016 <dbl>, YR2017 <dbl>, YR2018 <dbl>
```

- `contains()`: Contains a literal string

Example Pick all those variables containing the word `color` from the `starwars` dataset.

```
starwars %>% select(name, contains("color"))
```

```
## # A tibble: 87 x 4
##   name      hair_color  skin_color  eye_color
##   <chr>      <chr>      <chr>      <chr>
```

```
## 1 Luke Skywalker      blond      fair      blue
## 2 C-3PO                <NA>       gold      yellow
## 3 R2-D2                <NA>       white, blue red
## 4 Darth Vader         none        white      yellow
## 5 Leia Organa         brown       light      brown
## 6 Owen Lars           brown, grey light      blue
## 7 Beru Whitesun lars brown       light      blue
## 8 R5-D4               <NA>       white, red red
## 9 Biggs Darklighter  black       light      brown
## 10 Obi-Wan Kenobi     auburn, white fair      blue-gray
## # ... with 77 more rows
```

- `num_range()`: Matches a numerical range like x01, x02, x03

Example Select GDP data and country names from `g3` during 2005 and 2010.

```
g3 %>% select(`Country Name`, num_range("YR", 2005:2010))
```

```
## # A tibble: 264 x 7
##   `Country Name`      YR2005 YR2006 YR2007 YR2008 YR2009 YR2010
##   <chr>              <dbl>  <dbl>  <dbl>  <dbl>  <dbl>  <dbl>
## 1 Afghanistan      11.2    5.36 13.8    3.92   21.4   14.4
## 2 Albania           5.53    5.90 5.98    7.50    3.35   3.71
## 3 Algeria           5.91    1.68 3.37    2.36    1.63   3.63
## 4 American Samoa   -0.402  -4.17 1.96   -2.61   -4.24   0.442
## 5 Andorra           7.40    4.54 0.0400 -8.59   -3.69  -5.36
## 6 Angola            15.0    11.5 14.0    11.2    0.859  4.86
## 7 Antigua and Barbuda 6.41   12.7  9.26   -0.0301 -12.1  -7.20
## 8 Arab World        5.72    6.50 4.57    5.82    0.428  4.77
## 9 Argentina         8.85    8.05 9.01    4.06   -5.92  10.1
## 10 Armenia          13.9    13.2 13.7    6.90   -14.1   2.20
## # ... with 254 more rows
```

- `ends_with()`: Ends with a suffix
- `one_of()`: Matches variable names in a character vector
- `everything()`: Matches all variables
- `last_col()`: Select last variable, possibly with an offset
- `matches()`: Matches a regular expression (a sequence of symbols/characters expressing a string/pattern to be searched for within text)

relocate()

`relocate()` helps you organize columns by changing column positions.

Example Take `beatles.catalog`. Reorder columns such that year appears first.

```
beatles.catalog %>% relocate(year)
```

```
## # A tibble: 3 x 3
##   year album          num.tracks
##   <dbl> <chr>          <dbl>
## 1 1963 Please Please Me          14
## 2 1965 Rubber Soul             14
## 3 1967 Magical Mystery Tour      11
```

You can also reorder columns by their types. For example, if you wish to organize `beatles.catalog` such that numeric objects appear first followed by character, you can do this by writing: `beatles.catalog %>% relocate(where(is.numeric))`

`summarize()`

`summarize()` summarizes variables—you choose the variables and the summaries (e.g., `mean()` or `min()`).

```
df %>% summarize(
  mean(x), mean(y),
  min(x), max(x),
  min(y), max(y)
)
```

```
##   mean(x) mean(y) min(x) max(x) min(y) max(y)
## 1      12      11      2     22      1     20
```

returns a 1×6 tibble with the means of `x`, `y`; the minimum of `x` and `y`; and the maximum of `x` and `y`.

`summarize()` and `group_by()`

`group_by()` groups your observations by the variable(s) that you name.

Specifically, `group_by()` returns a *grouped data frame* that you can then feed to `summarize()`, `mutate()` to perform grouped calculations, e.g., each group's mean.

Example: Grouped summaries

```
# Create a new data frame
our_df <- tibble(
  df,
  grp = rep(c("A", "B"), each = 3)
)
```

```
## # A tibble: 6 x 3
##       x     y grp
##   <dbl> <int> <chr>
## 1     2    20  A
## 2     6     1  A
## 3    10    10  A
## 4    14     5  B
## 5    18    14  B
## 6    22    16  B
```

```
# For dataset 'our_df' ...
our_df %>%
  # Group by 'grp'
  group_by(grp) %>%
  # Take means of 'x' and 'y'
  summarize(mean(x), mean(y))
```

```
##   mean(x) mean(y)
## 1      12      11
```


Example: Grouped mutation

```
# Create a new data frame
our_df <- data.frame(
  df,
  grp = rep(c("A", "B"), each = 3)
)

##    x  y grp
## 1  2 20  A
## 2  6  1  A
## 3 10 10  A
## 4 14  5  B
## 5 18 14  B
## 6 22 16  B
```

```
# Add grp means for x and y
our_df %>%
  group_by(grp) %>%
  mutate(
    x_m = mean(x), y_m = mean(y)
  )

## # A tibble: 6 x 5
## # Groups:   grp [2]
##       x     y grp   x_m  y_m
##   <dbl> <int> <chr> <dbl> <dbl>
## 1     2    20  A      12    11
## 2     6     1  A      12    11
## 3    10    10  A      12    11
## 4    14     5  B      12    11
## 5    18    14  B      12    11
## 6    22    16  B      12    11
```

filter()

The filter() function does what its name implies: it **filters the rows** of your data frame **based upon logical conditions**.

Example:

```
# Create a dataset
some_df <- data.frame(
  x = 1:10,
  y = 11:20
)
```

```
# Only keep rows where x is 3
some_df %>% filter(x == 3)

##    x  y
## 1  3 13
```

Using the same dataset and filter, perform the following operations-

- keep rows where $x \geq 6$
- keep rows where $y/x \geq 2$
- keep rows where $12 \leq y \leq 18$

arrange()

arrange() will sort the rows of a data frame using the inputted columns.

R defaults to starting with the “lowest” (smallest) at the top of the data frame. Use a - in front of the variable’s name to reverse sort.

```
# As is
```

```
our_df
```

```
##      x y grp
## 1   2 20  A
## 2   6  1  A
## 3  10 10  A
## 4  14  5  B
## 5  18 14  B
## 6  22 16  B
```

```
# Arrang by y, grp, then -x
```

```
our_df %>% arrange(y, grp, -x)
```

```
##      x y grp
## 1   6  1  A
## 2  14  5  B
## 3  10 10  A
## 4  18 14  B
## 5  22 16  B
## 6   2 20  A
```

slice()

arrange() will subset the data frame using the row index provided by you.

```
n_rows <- 12:18
```

```
slice(mtcars, n_rows)
```

```
##      mpg cyl  disp  hp drat   wt  qsec vs am gear carb
## Merc 450SE    16.4   8  275.8 180 3.07 4.070 17.40  0  0    3    3
## Merc 450SL    17.3   8  275.8 180 3.07 3.730 17.60  0  0    3    3
## Merc 450SLC    15.2   8  275.8 180 3.07 3.780 18.00  0  0    3    3
## Cadillac Fleetwood 10.4   8  472.0 205 2.93 5.250 17.98  0  0    3    4
## Lincoln Continental 10.4   8  460.0 215 3.00 5.424 17.82  0  0    3    4
## Chrysler Imperial 14.7   8  440.0 230 3.23 5.345 17.42  0  0    3    4
## Fiat 128      32.4   4   78.7  66 4.08 2.200 19.47  1  1    4    1
```

distinct()

distinct() will remove duplicates from your data.

```
name <- c("Siddharth", "Rajshree", "Ankitha", "Ankitha")
```

```
CGPA <- c(3.9, 3.5, 3.4, 3.4)
```

```
age <- c(21, 22, 24, 24)
```

```
school.db <- data.frame(name, CGPA, age)
```

```
print(school.db)
```

```
##      name CGPA age
## 1 Siddharth 3.9  21
## 2 Rajshree 3.5  22
## 3 Ankitha 3.4  24
## 4 Ankitha 3.4  24
```

```
school.db %>% distinct()
```

```
##      name CGPA age
## 1 Siddharth 3.9  21
## 2 Rajshree 3.5  22
## 3 Ankitha 3.4  24
```

Chain Operation Revisited

Let's combine several dplyr operations into a chain.

Select data → Select groups → Select columns → Compute averages for selected columns

In this example, we will calculate average height and mass by species and sex.

```

starwars %>%                                #select data
  group_by(species, sex) %>%                 #group variables
  select(height, mass) %>%                 #select columns
  dplyr::summarise(                         #compute averages
    ave.height = mean(height, na.rm = TRUE),
    ave.mass = mean(mass, na.rm = TRUE)
  )

```

```
## Adding missing grouping variables: `species`, `sex`
```

```
## `summarise()` regrouping output by 'species' (override with `.groups` argument)
```

```
## # A tibble: 41 x 4
```

```
## # Groups:   species [38]
```

```

##   species sex    ave.height ave.mass
##   <chr>   <chr>      <dbl>   <dbl>
## 1 Aleena  male         79        15
## 2 Besalisk male        198       102
## 3 Cerean  male        198        82
## 4 Chagrian male        196       NaN
## 5 Clawdite female      168        55
## 6 Droid   none        131.       69.8
## 7 Dug     male        112        40
## 8 Ewok    male         88        20
## 9 Geonosian male       183        80
## 10 Gungan male       209.        74
## # ... with 31 more rows

```

Done for the day

```
## Sorry, this silly GIF is only available in the the HTML version of the notes.
```